

Raymond Yang

raymondmyang@gmail.com | 2069736678 | Seattle, WA

EDUCATION BACKGROUND

University of Washington

Bachelor of Science in Data Science GPA 3.5/4.0

Seattle, WA

September 2021 - March 2025

TECHNICAL SKILLS

- **Analytical Skills:** Data visualization, Reporting Analysis, Data Cleaning, KPI tracking, Statistical Analysis, Business Acumen, Data Modeling, Forecasting Analysis, ETL/Data Pipeline, Data Governance, Data Mining, Database Querying, AI/ML
- **Programming Languages:** SQL, Python, R, SAS, VBA, JavaScript, Matlab, Bash/shell scripting, DAX, Java, C, Html
- **Tools & Platforms:** Advanced Excel(Pivot Tables, Power Query, VBA, INDEX/VLOOKUP), Tableau, Power BI, SQL Database, Looker, Snowflake, Google BigQuery, AWS, Azure, dbt, Alteryx, SPSS, Jira, Github, Google Analytics, Smartsheet, Apache Airflow, Jupyter Notebook, Oracle Cloud/Fusion, SharePoint, Power Platform, Databricks, ChromaDB, Ollama
- **Libraries:** Pandas, NumPy, Matplotlib, SciPy, Seaborn, PySpark, Scikit-learn, dplyr/tidyverse, Tensorflow, LangChain, LangGraph
- **Soft Skills:** Stakeholder Communication, Problem Solving, Critical Thinking, Attention to detail, Presentation Skill, Quick Learner, Collaboration/Team Work, Data StoryTelling, Analytical Thinking, Stakeholder Management

WORK EXPERIENCE

Allen Institute

Data Scientist

Seattle, WA

Jan 2025 – Dec 2025

- Designed and implemented 20+ Smartsheet reports and Power BI dashboards for 6 Scientific Departments, implemented formula calculations from 10+ Excel sheets, dealt with over 100 million\$ funding and 198 grants.
- Built an analytics data mart for grants and funding in Snowflake, with ingestion & transformation jobs orchestrated from Databricks notebooks and Python scripting; documented models and lineage in Git repos and automated CI checks.
- Wrote SQL for quality checks and data validation across Smartsheet/Oracle sources; added regression-style predictive analysis (Scikit-learn/XGBoost) to forecast run-rate and grant burn, with results shared via Jupyter/Colab notebooks.
- Partnered with PIs, grants administrators, and department leadership across 6 scientific departments to translate business requirements into analytics deliverables; reducing inbound reporting requests from department staff by ~30%.

Codeflag AI(Shanghai)Co., Ltd

Data Engineer

Shanghai, China

May 2024 – Sept 2024

- Architected an automated Azure medallion ETL pipeline (ADF + Data Lake Gen2 + Databricks + Synapse Analytics), reducing manual reporting prep time by ~70% and delivering daily fresh KPI data to stakeholders.
- Engineered dynamic ForEach-driven ADF pipelines that auto-discovered all SQL schema tables and triggered downstream PySpark transformation jobs across layers,eliminating hardcoded references and cutting pipeline maintenance overhead by ~4 hours per sprint.
- Enforced RBAC and secret management via Azure Key Vault and Microsoft Entra ID, securing SQL Server credentials and Databricks access tokens for all pipeline service identities and eliminating plaintext credential exposure across a 3-person team of data engineers and cloud architect.

PROJECT EXPERIENCE

Vantage AI Research Tool — Financial Intelligence

Seattle, WA

AI/ML Engineer

Mar 2025 – Jun 2025

- Built a local RAG agent (Ollama + ChromaDB + LangChain) for CSV document semantic retrieval and a LangGraph research agent that simultaneously scraped Google, Bing, and Reddit via Bright Data APIs, synthesizing multi-source findings through GPT-4o.
- Leveraged LangGraph and LangChain as the core orchestration layer with Ollama(llama3.2) for local inference and GPT-4o for cloud inference; Pydantic for structured LLM outputs, and Bright Data SERP/Social APIs for live web data ingestion.
- Collaborated with data scientist, backend engineers and product leads to define LangGraph state schema and implement snapshot polling retry logic; led weekly prompt iteration sessions standardizing system/user prompt pairs across 5 analysis stages.
- Eliminated \$2,800/year in cloud LLM API spend via on-premise Ollama inference; the multi-source research agent cut analyst research cycle time by 62%, enabling the team to handle 3× more client intelligence requests per sprint without adding headcount.

Serverless Web Application & Analytics Pipeline

Seattle, WA

AWS Engineer/Software Engineer

Mar 2024 – Jun 2024

- Architected a serverless web application on AWS integrating 7 cloud services (Amplify, Cognito, Lambda, API Gateway, DynamoDB, IAM, GitHub), establishing a CI/CD pipeline that reduced manual deployment overhead by ~70% and enabled continuous delivery to production.
- Designed a secure RESTful API layer using Amazon API Gateway with Cognito JWT-based authorization, backed by a Node.js Lambda function performing real-time ride dispatch and DynamoDB writes, eliminating EC2 infrastructure costs.
- Engineered an end-to-end AWS analytics pipeline integrating AWS services (DynamoDB, Lambda, Glue, Athena, QuickSight), translating NoSQL data into SQL-queryable form and eliminating manual ETL exports to enable self-serve real-time reporting.